

Interpreting the coefficients in Linear Models

This section applies only to Linear models (`LinearRegression` ,
`LogisticRegression`)

It may seem overly specialized, but since these models are used so often, we will spend some time.

Also, because we can assign a meaning to the coefficients, these models are highly interpretable.

Numeric features

For Linear Regression

$$\hat{\mathbf{y}} = \Theta^T \cdot \mathbf{x} = \sum_{j=1}^n \Theta_j * \mathbf{x}_j$$

so for a *unit* change in \mathbf{x}_j :

$$\Delta \hat{\mathbf{y}} = \Theta_j$$

$$\text{Thus } \Theta_j = \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{x}_j}$$

That is, the coefficient Θ_j is the amount $\hat{\mathbf{y}}$ changes for a 1 unit change in \mathbf{x}_j

Does this mean that, for two features j, j' if

- $\Theta_j > \Theta_{j'}$
- that feature j is more important than j' ?

No.

Unless \mathbf{x}_j and $\mathbf{x}_{j'}$ are on same scale (e.g., have been standardized) a change of 1 unit isn't comparable.

For example

- \mathbf{x}_j in miles, $\mathbf{x}_{j'}$ in inches

What Θ_j does tell you is

- the direction
- and magnitude

of how much the prediction changes when the feature changes

This can be useful for interpretation

- In our housing price premium example: premium *increases* with area since $\Theta_j > 0$

Remember that transformations change units of the variables (features or targets).

So the Θ_j reflects the sensitivity to the change in the *transformed* units, not the original.

Recall the target for Logistic Regression was transformed to log odds

$$y = \log\left(\frac{\hat{p}}{1 - \hat{p}}\right)$$

So a unit change in feature \mathbf{x}_j with parameter Θ_j changes the *odds* $\frac{\hat{p}}{1 - \hat{p}}$ in a *multiplicative* way

$$\log\left(\frac{\hat{p}}{1 - \hat{p}}\right) + \Theta_j = \log\left(\frac{\hat{p}}{1 - \hat{p}} * \exp \Theta_j\right)$$

-

Examples

- Log transform of target:
 - $\log y = \Theta_0 + \Theta_1 * x_1$
 - $\theta_1 = \frac{\partial \log y}{\partial x_1} = \% \text{ change in } y \text{ per unit change in } x_1$
- Log transform of both target and feature:
 - $\log y = \Theta_0 + \Theta_1 * \log x_1$
 - $\Theta_1 = \frac{\partial \log y}{\partial \log x_1} = \% \text{ change in } y \text{ per } \% \text{ change in } x_1$
- Standardize feature
 - Transform x into $z_x = \frac{x - \bar{x}}{\sigma_x}$
 - $y = \Theta_0 + \Theta_1 * z_x$
 - $\Theta_1 = \frac{\partial \log y}{\partial z_x}$ change in y per 1 standard deviation change in x
 - since z is in units of "number of standard deviations"

Remember

- if you transform features in training, you must apply the same transformation to features in test
 - if the transformation is parameterized, the parameters are determined at **train** fit time, not test !
- if you transform the target, the prediction is in different units than the original
 - you can perform the inverse transformation to get a prediction in original units

Categorical features

Consider when x_j is a binary categorical feature: $x_j \in \{0, 1\}$

Then Θ_j is the increase in \mathbf{y} associated with \mathbf{x}_j being equal to 1 rather than 0.

Just like with numeric features.

Recall our discussion about the dummy variable trap for Linear Regression:

- if we have an intercept (with parameter Θ_0)
 - then Θ_0 is the *contribution* (not increment) to \mathbf{y} when $\mathbf{x}_j = 0$
 - as you can see from equation $\mathbf{y} = \Theta_0 + \dots + \Theta_j * \mathbf{x}_j$, when $\mathbf{x}_j = 0$

What's wrong with representing multinomial categorical values as numbers ?

We have already given several reasons why representing Categorical features as numeric is a bad idea.

Now that we can interpret Θ_j , we add another reason to the list.

Let's consider the Passenger Class (*PClass*) variable from the Titanic example:

$PClass \in \{1, 2, 3\}$

- The difference in prediction for
 - (*PClass* = 1) vs (*PClass* = 2),
 - or (*PClass* = 2) vs (*PClass* = 3)
 - is Θ_{PClass}
- BUT the difference in prediction for *PClass* = 1 vs (*PClass* = 3) is $2 \times \Theta_{PClass}$
 - Is y (log odds of Survival) impacted by a factor of 2 ?
 - That's what you are saying by representing *PClass* as the integer range [1, 3]

- What if $P_{\text{class}} \in \{100, 200, 300\}$?
 - Does changing one class impact log odds of Survival by $100 * \Theta_{P_{\text{Class}}}$?

Categorical variables have neither order nor magnitude; integers have both.

Beware of representing Categorical variables with numbers.

Bucketing/Binning re-visited

Suppose \mathbf{x}_j is a continuous numeric feature (e.g., Age).

Then a 1 year change in Age means that \mathbf{y} changes by Θ_j .

Do you think that 1 year of increase is as important for an adult as for an infant ?

- That's what the regression equation is telling you
- If not, consider
 - Bucketing/Binning continuous variables
 - Percent changes

Interpreting the MNIST classifier: template matching

The Θ produced by a linear classifier can be viewed as templates

- the strength of Θ_j tells you how strongly feature \mathbf{x}_j influences the target

So we can interpret Θ as a "template" for what a model is looking for.

Let's look at the template for

- The 10 separate, single-digit binary MNIST classifiers
- Or similarly: each row of Θ for the multinomial 10 class MNIST classifier

Here's the training data

```
In [5]: mnh.setup()  
        mnh.visualize()
```

Retrieving MNIST_784 from cache



Let's fit a `LogisticRegression` model and examine the templates (coefficients Θ)


```
In [6]: mnh.fit()  
mnist_fig, mnist_ax = mnh.plot_coeff()
```

Example run in 3.128 s


```
Out[6]: LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,  
    intercept_scaling=1, max_iter=100, multi_class='multinomial',  
    n_jobs=None, penalty='l2', random_state=None, solver='saga',  
    tol=0.1, verbose=0, warm_start=False)
```

```
/home/kjp/anaconda3/lib/python3.7/site-packages/matplotlib/figure.py:445: User  
Warning: Matplotlib is currently using module://ipykernel.pylab.backend_inlin  
e, which is a non-GUI backend, so cannot show the figure.  
  % get_backend())
```

Parameters for...



Class 0




Class 1



Class 2



Class 3



Class 4



Class 5



Class 6



Class 7



Class 8



Class 9

Recall

- There is one parameter per pixel
- The parameters are ordered in the same way as the linearization of the pixels
 - from (28×28) grid to a vector of 784 numbers.
- We can display the 784 parameters in a (28×28) image to show the intensity of parameter associated with a pixel
- White is high parameter value; Black is low (or negative)

- The template for 0 emphasizes small values (absence of bright pixels) in the center of the image
- The template for 1 emphasizes bright vertical pixels
- The template for 8 emphasizes the absence of bright pixels
 - in the two circles
 - in the pinched waist

You can now imagine how these templates might lead to misclassification

What is the classification of

- a "7" with a strong vertical line in the center (that's what the "1" template tries to match)
- a thin "0" (the "0" template is looking for a large donut)

So interpretation is a very powerful diagnostic tool for both understanding and improving your models.

In [7]: `print("Done")`

Done