

Classification via counting

Our goal in classification is to construct the conditional probability distribution $p(\mathbf{y}|\mathbf{x})$

- given feature vector \mathbf{x} , \mathbf{y} is a vector of probabilities
- y_c is the probability that \mathbf{x} is in class c
 - **Notation abuse alert:** subscripts like c should be integers rather than class names
 - so technically should write $y_{m(c)}$ where $m(c)$ is the integer index of class named c

This sounds difficult at first glance.

Let's start with something simpler: counting.

We will show to to construct this probability using nothing more than counting the features and targets of the training set!

From counting to probability

We introduce the topic by assuming all our variables (features and target) are discrete.

We will subsequently adapt this to continuous variables.

First, let's compute the distribution of target classes.

Let $\mathbf{X} = \{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)}) | 1 \leq i \leq m\}$ be our m training examples.

Then

$$\text{\textcolor{red}{cnt}}_{\mathbf{y} = \mathbf{y}'} = |\{i \mid \mathbf{y}^{(i)} = \mathbf{y}'\}|$$

is the number of training examples with target \mathbf{y}' .

We can easily convert this into an unconditional probability

$$p(\mathbf{y} = \mathbf{y}') = \frac{\text{\textcolor{red}{cnt}}_{\mathbf{y} = \mathbf{y}'}}{m}$$

We can similarly compute the *joint* probability of any two features.

First we count the cooccurences of the two variables

$$\text{\textcolor{red}{cnt}}_{\mathbf{x}_j = \mathbf{x}'_j, \mathbf{x}_k = \mathbf{x}'_k} = \left| \{ i | \mathbf{x}_j^{(i)} = \mathbf{x}'_j, \mathbf{x}_k^{(i)} = \mathbf{x}'_k \} \right|$$

And the the joint probability is

$$p(\mathbf{x}_j = \mathbf{x}'_j, \mathbf{x}_k = \mathbf{x}'_k) = \frac{\text{\textcolor{red}{cnt}}_{\mathbf{x}_j = \mathbf{x}'_j, \mathbf{x}_k = \mathbf{x}'_k}}{m}$$

Our illustration is with two features but the notation generalizes for

- any number of variables
- any kind of variables: feature or target

Finally, we can define *conditional* probability

$$p(\mathbf{y} = \mathbf{y}' | \mathbf{x} = \mathbf{x}') = \frac{p(\mathbf{y} = \mathbf{y}', \mathbf{x} = \mathbf{x}')}{p(\mathbf{x} = \mathbf{x}')}$$

That is, the conditional probability

- is the joint probability
- as a fraction, relative to the unconditional probability of $\mathbf{x} = \mathbf{x}'$

Bayes theorem

The key for converting counts (really, associated probabilities) to predictions lies in Bayes Theorem.

Bayes Theorem relates conditional and unconditional probabilities.

Bayes Theorem

$$p(\mathbf{y} = \mathbf{y}' | \mathbf{x} = \mathbf{x}') = \frac{p(\mathbf{x} = \mathbf{x}' | \mathbf{y} = \mathbf{y}') * p(\mathbf{y} = \mathbf{y}')}{p(\mathbf{x} = \mathbf{x}')}$$

Let's think about Bayes Theorem in terms of our classification task:

- The left hand side is our prediction for the class probabilities, given the features
- The right hand side involves
 - the conditional probability of seeing examples with features \mathbf{x}' and target \mathbf{y}' .
 - the unconditional probability of seeing examples with label \mathbf{y}
 - the unconditional probability of seeing examples with feature vector \mathbf{x} .

All these elements can be obtained by counting (and filtering) the training set !

Hence, we can build an extremely simple classifier using nothing more than counting.

Posterior, Prior Probability, Evidence

Let's break down the parts of Bayes theorem and give them some names:

- $p(\mathbf{y} = \mathbf{y}' | \mathbf{x} = \mathbf{x})$: posterior probability
 - our prediction
 - this is the probability distribution of \mathbf{y} *conditional* on the features being \mathbf{x}
- $p(\mathbf{y} = \mathbf{y}')$: prior probability
 - this is the unconditional distribution of \mathbf{y}
- $p(\mathbf{x} = \mathbf{x}' | \mathbf{y} = \mathbf{y})$: likelihood
 - given that $\mathbf{y} = \mathbf{y}'$, what is the probability that $\mathbf{x} = \mathbf{x}'$?
 - this is the counting part: how often does the label \mathbf{y} occur when the features are \mathbf{x} ?
- $p(\mathbf{x} = \mathbf{x}')$: evidence
 - how often do we see the features \mathbf{x} ?

We can re-state Bayes Theorem as $\text{posterior} = \frac{\text{prior} * \text{likelihood}}{\text{evidence}}$

That is:

- starting from an uninformed *prior* distribution of \mathbf{y}
- derive a conditional *posterior* distribution (i.e., informed by *evidence* \mathbf{x}) by updating via the *likelihood* of seeing \mathbf{x}, \mathbf{y} together.

Proof of Bayes Theorem

$$p(\mathbf{y} = \mathbf{y}' | \mathbf{x} = \mathbf{x}') = \frac{p(\mathbf{y}=\mathbf{y}', \mathbf{x}=\mathbf{x}')}{p(\mathbf{x}=\mathbf{x}')} \quad (\text{def. of conditional probability})$$

$$= \frac{p(\mathbf{y}=\mathbf{y}', \mathbf{x}=\mathbf{x}')}{p(\mathbf{x}=\mathbf{x}')} * \frac{\frac{1}{p(\mathbf{y}=\mathbf{y}')}}{\frac{1}{p(\mathbf{y}=\mathbf{y}')}} \quad (\text{multiply by identity})$$

$$= \frac{p(\mathbf{x}=\mathbf{x}' | \mathbf{y}=\mathbf{y}')}{p(\mathbf{x}=\mathbf{x}')} * \frac{1}{\frac{1}{p(\mathbf{y}=\mathbf{y}')}} \quad (\text{def. of conditional probability})$$

$$= \frac{p(\mathbf{x}=\mathbf{x}' | \mathbf{y}=\mathbf{y}')}{p(\mathbf{x}=\mathbf{x}')} * p(\mathbf{y} = \mathbf{y}')$$

Length of \mathbf{x} is n

Remember that \mathbf{x} is a vector, so that $p(\mathbf{x} = \mathbf{x}' | \mathbf{y} = \mathbf{y}')$ is a *joint* probability of n terms

$$p(\mathbf{x}_1 = \mathbf{x}'_1, \mathbf{x}_2 = \mathbf{x}'_2, \dots, \mathbf{x}_n = \mathbf{x}'_n | \mathbf{y} = \mathbf{y}')$$

We can obtain this by counting (as described above)

- Let cnt_{x_j} denote the number of distinct values for the j^{th} feature
- We need to count

$$\prod_{1 \leq j \leq n} \text{cnt}_{x_j}$$

combinations

That's a lot of counting. More importantly, it's a lot of parameters to remember (i.e, size of Θ is big).

We need a short-cut.

The Naive part of Naive Bayes

We will assume that each feature is *conditionally* independent of one another

$$p(\mathbf{x}_j = \mathbf{x}'_j, \mathbf{x}_k = \mathbf{x}'_k, |\mathbf{y} = \mathbf{y}') = p(\mathbf{x}_j = \mathbf{x}'_j | \mathbf{y} = \mathbf{y}') * p(\mathbf{x}_k = \mathbf{x}'_k | \mathbf{y} = \mathbf{y}')$$

That is

- x_j and x_k are *not* independent
- they are independent *conditional* on $\mathbf{y} = \mathbf{y}'$

Think of x_j and x_k being correlated through their individual relationships with \mathbf{y} .

Excluding that mutual dependence, they may be uncorrelated.

if we make the "naive" assumption of conditional independence

$$p(\mathbf{x} = \mathbf{x}' | \mathbf{y} = \mathbf{y}') = \prod_{i=1}^n p(\mathbf{x}_i = \mathbf{x}'_i | \mathbf{y} = \mathbf{y}')$$

That is, the joint conditional probability of the vector of length n is **assumed** to be the product of the individual conditional probabilities of each element of the vector.

This is probably not true, but

- is very easy to compute
- uses few parameters
- may be close enough

What about computing the *unconditional* $p(\mathbf{x} = \mathbf{x}')$?

We can obtain this from conditional probabilities as well

$$p(\mathbf{x} = \mathbf{x}') = \sum_{c \in C} p(\mathbf{x} = \mathbf{x}' | \mathbf{y} = c) * p(\mathbf{y} = c)$$

That is, the unconditional probability follows from the

- conditional probability given \mathbf{y}
- weighted by the probability $p(\mathbf{y})$ for each possible value of \mathbf{y}

This follows from the definition of conditional probability.

What this means is that the only parameters we need to remember are

- the unconditional probabilities $p(\mathbf{y})$
 - depends on number of classes $|C|$
- probabilities conditional on \mathbf{y} : $p(\mathbf{x}|\mathbf{y})$
 - depends on length of \mathbf{x} : n

Example

Here is a hypothetical trading example for equities

- there are two attributes
 - Valuation: Rich/Cheap
 - is the current stock price expensive (Rich) or inexpensive (Cheap) ?
 - Yield: High/Low
 - is the dividend yield of the stock desirable (High) or undesirable (Low) ?
- target: Long/Short/Neutral
 - What should our position be ?

We are given a number of examples and want a classifier to tell us

- for a new equity
 - given it's Valuation and Yield: what should our position be ?

Here are our observations

```
In [4]: d_df = pd.read_csv("valuation_yield_action.csv")
target_name = "Action"
d_df
```

Out[4]:

	Valuation	Yield	Action
0	Cheap	High	Long
1	Cheap	High	Long
2	Cheap	High	Long
3	Cheap	High	Neutral
4	Rich	Low	Short
5	Rich	Low	Short
6	Rich	Low	Short
7	Rich	Low	Short
8	Rich	Low	Neutral
9	Cheap	Low	Neutral
10	Cheap	Low	Long
11	Cheap	Low	Short
12	Cheap	Low	Neutral
13	Rich	High	Long
14	Rich	High	Short
15	Rich	High	Neutral
16	Fair	Low	Long
17	Fair	Low	Short
18	Fair	Low	Neutral
19	Fair	High	Long
20	Fair	High	Short
21	Fair	High	Neutral
22	Cheap	Low	Long
23	Fair	Low	Short
24	Fair	High	Long

And a quick look at the data, slice by Action

```
In [5]: grouped_by_target = d_df.groupby(target_name)
        for gp in grouped_by_target.groups.keys():
            print(gp, "\n")
            print(grouped_by_target.get_group(gp).head())
```

Long

	Valuation	Yield	Action
0	Cheap	High	Long
1	Cheap	High	Long
2	Cheap	High	Long
10	Cheap	Low	Long
13	Rich	High	Long

Neutral

	Valuation	Yield	Action
3	Cheap	High	Neutral
8	Rich	Low	Neutral
9	Cheap	Low	Neutral
12	Cheap	Low	Neutral
15	Rich	High	Neutral

Short

	Valuation	Yield	Action
4	Rich	Low	Short
5	Rich	Low	Short
6	Rich	Low	Short
7	Rich	Low	Short
11	Cheap	Low	Short

Looks like we

- go Long if the stock is Cheap (valuation) and High (Yield)
- go short if the stock is Rich (expensive valuation) and Low (Yield)

Here's the empirical distribution of the training examples

```
In [6]: d_df["dummy"] = 1 # Need to aggregate on something
t = d_df.pivot_table(index=target_name, columns=["Valuation", "Yield"], values=
"dummy", aggfunc=["count"],
                    margins=True)

t
```

Out[6]:

	count						
Valuation	Cheap		Fair		Rich		All
Yield	High	Low	High	Low	High	Low	
Action							
Long	3.0	2.0	2.0	1.0	1.0	NaN	9
Neutral	1.0	2.0	1.0	1.0	1.0	1.0	7
Short	NaN	1.0	1.0	2.0	1.0	4.0	9
All	4.0	5.0	4.0	4.0	3.0	5.0	25

This gives us everything we need for the Naive Bayes algorithm

- these are counts
- we can easily turn the counts into unconditional probabilities by dividing by total number of examples
- will leave them as counts for now

Let's parse this table:

- Columns: $\text{cnt}_{y|\mathbf{x}}$
 - a column (defined by concrete values for each of the two attributes)
 - defines a distribution over the target (Action)
- Column Sum: $\text{cnt}_{\mathbf{x}} = \sum_{a \in \text{Action}} \text{cnt}_{a|\mathbf{x}}$
 - total number of examples with attribute pair \mathbf{x}
- Rows: $\text{cnt}_{\mathbf{x}|y}$
 - a row (defined by a concrete value for the Action)
 - defines a distribution over the attributes pairs for which this action is taken
- Row sums: $\text{cnt}_a = \sum_{\mathbf{x}} \text{cnt}_{\mathbf{x}|a}$
 - total number of examples with Action a

Let's simplify the table by looking at the marginal with respect to each attribute

- distribution over a single attribute rather than the pair

First, by Valuation

```
In [7]: # Single feature (Valuation), rather than pair
d_df.drop(columns=["dummy"]).pivot_table(index=target_name, columns=["Valuation"
], aggfunc=["count"],
fill_value=0, margins=True)
```

Out[7]:

count				
Yield				
Valuation	Cheap	Fair	Rich	All
Action				
Long	5	3	1	9
Neutral	3	2	2	7
Short	1	3	5	9
All	9	8	8	25

And by Yield

```
In [8]: # Single feature (Yield), rather than pair
d_df.drop(columns=["dummy"]).pivot_table(index=target_name, columns=["Yield"],
aggfunc=["count"],
fill_value=0, margins=True)
```

Out[8]:

	count		
	Valuation		
Yield	High	Low	All
Action			
Long	6	3	9
Neutral	3	4	7
Short	2	7	9
All	11	14	25

And the target (Action) distribution

In [9]: `t.loc[:, idx["count","All",:]]`

Out[9]:

count	
Valuation	All
Yield	
Action	
Long	9
Neutral	7
Short	9
All	25

Here is the target distribution as probabilities rather than counts


```
In [10]: num_examples = t.loc["All", idx["count", "All", :]][0]

print("There are {e:d} training examples".format(e=int(num_examples)) )

# Class probabilities
t.loc[:, idx["count", "All", :]]/t.loc["All", idx["count", "All", :]]
```

There are 25 training examples

Out[10]:

count	
Valuation	All
Yield	
Action	
Long	0.36
Neutral	0.28
Short	0.36
All	1.00

Why not just use the empirical distribution ?

At this point, it's fair to ask:

- given a test example $\mathbf{x} = \mathbf{x}'$
- we can read $p(\mathbf{y} = \mathbf{y}' | \mathbf{x} = \mathbf{x}')$ *directly from the table*

Why do we need Naive Bayes ?

Answer: Because the table can be big !

- one entry for every possible combination of features

The conditional independence assumption allows us

- to have a single vector for each feature $\mathbf{x}_1, \mathbf{x}_2$ individually: total $|\mathbf{x}_1| + |\mathbf{x}_2|$
- rather than *combinations* of features $\mathbf{x}_1, \mathbf{x}_2$: total $|\mathbf{x}_1| * |\mathbf{x}_2|$
 - in this case $n = 2$ but in general: total $\prod_{j=0}^n |\mathbf{x}_j|$

This is usually much smaller.

Drawbacks

The zero frequency problem

In order for Naive Bayes to work we must have

$$p(\mathbf{x}_j = \mathbf{x}'_j | \mathbf{y} = \mathbf{y}')$$

for *all* possible values of \mathbf{x}'_j that we will encounter during *inference* (test) time.

There is no guarantee that we will see each of these values in the training set.

- especially when the training set is small

If we don't, the probability is 0, which is not only probably wrong but can cause problems.

Additive smoothing

[additive smoothing](https://en.wikipedia.org/wiki/Additive_smoothing) (https://en.wikipedia.org/wiki/Additive_smoothing)

There is a simple solution to the zero frequency problem

- artificially inflate all counts by some parameter α .

This eliminates zero counts at the cost of biasing all counts.

Note that when converting counts to probabilities

- we have increased the count of each of the $|C|$ classes by α
- so the total count for the denominator is $m + |C| * \alpha$

Replace empirical distributions by functional forms

Another way to address the zero frequency problem is to avoid the empirical distribution of training data (the counts)

- assume the features come from a parameterized distribution
 - Bernoulli distribution for binary variables
 - Multinomial distribution for variables with more than two classes

This also has the advantage of fewer parameters: one parameter per feature.

Assumption of conditional independence

This is a questionable assumption.

In its defense:

- if n (the number of features) is very large
 - the conditional independence assumption is more likely to hold.

Advantages

- Very simple: just counting !
 - Easy and powerful Baseline Model to use in your Recipe for Machine Learning

Continuous variables for features

The above discussion was limited to features that could take on discrete values.

We now discuss how to include features that are continuous variables.

Discretizing continuous variables

The simplest way to deal with a continuous feature \mathbf{x}_j is to turn it into one or more discrete variables.

- define a threshold t_j and replace the continuous \mathbf{x}_j with a binary variable
 - $I_{\mathbf{x}_j < t_j}$
- define multiple intervals on the range of \mathbf{x}_j and create a binary variable per interval
 - $I_{t_{j,l-1} \leq x_j \leq t_{j,l}}$
- the thresholds are a hyper-parameter: can search for optimal

Unfortunately the ordering relationship between continuous values is lost

- we have made them categorical

We see this same technique used in Decision Trees, so it's worth mentioning.

Replacing empirical distributions by functional forms for continuous variables

Another technique for continuous variables

- Replace the discrete empirical distribution by a functional form
 - Gaussian

This has the advantage that many distributions are characterized by a small number of parameters

- Gaussian: 2 per feature -- a mean and standard deviation

This also deals with the zero frequency problem by eliminating the empirical distribution.

```
In [11]: print("Done")
```

Done